

# MULTI THREADED PATTERN SEARCHING OF LARGE FILES USING LIMITED MEMORY

Submission by: Peter Morley  
Graduate Professor: Resit Sendag Ph.D.

# Pattern Discovery Approach

- Discover patterns using the suffix tree algorithm
- A single threaded pattern finding program exists but pattern searching speeds could be greatly improved by implementing a multithreaded algorithm
- Files that are larger than a computer's DRAM capacity must be partially processed using the Hard Disk otherwise processing can not be accomplished.

# Previous work on Pattern Discovery

- The “An Analysis of Address and Branch Patterns with PatternFinder” paper implemented a single threaded, branch and address specific pattern discovery tool.
  - Specific to branch and address discovery, not generic
  - Slow because it is single threaded
- The “Better External Memory Suffix Array Construction” paper discussed an implementation of the Suffix Tree algorithm when DRAM alone is not capable of processing large files.
  - Many unique algorithms but no mention of HD processing speed improvement techniques
  - Pattern Finder’s memory mapping solution improves HD read/write time

# Suffix Tree Algorithm

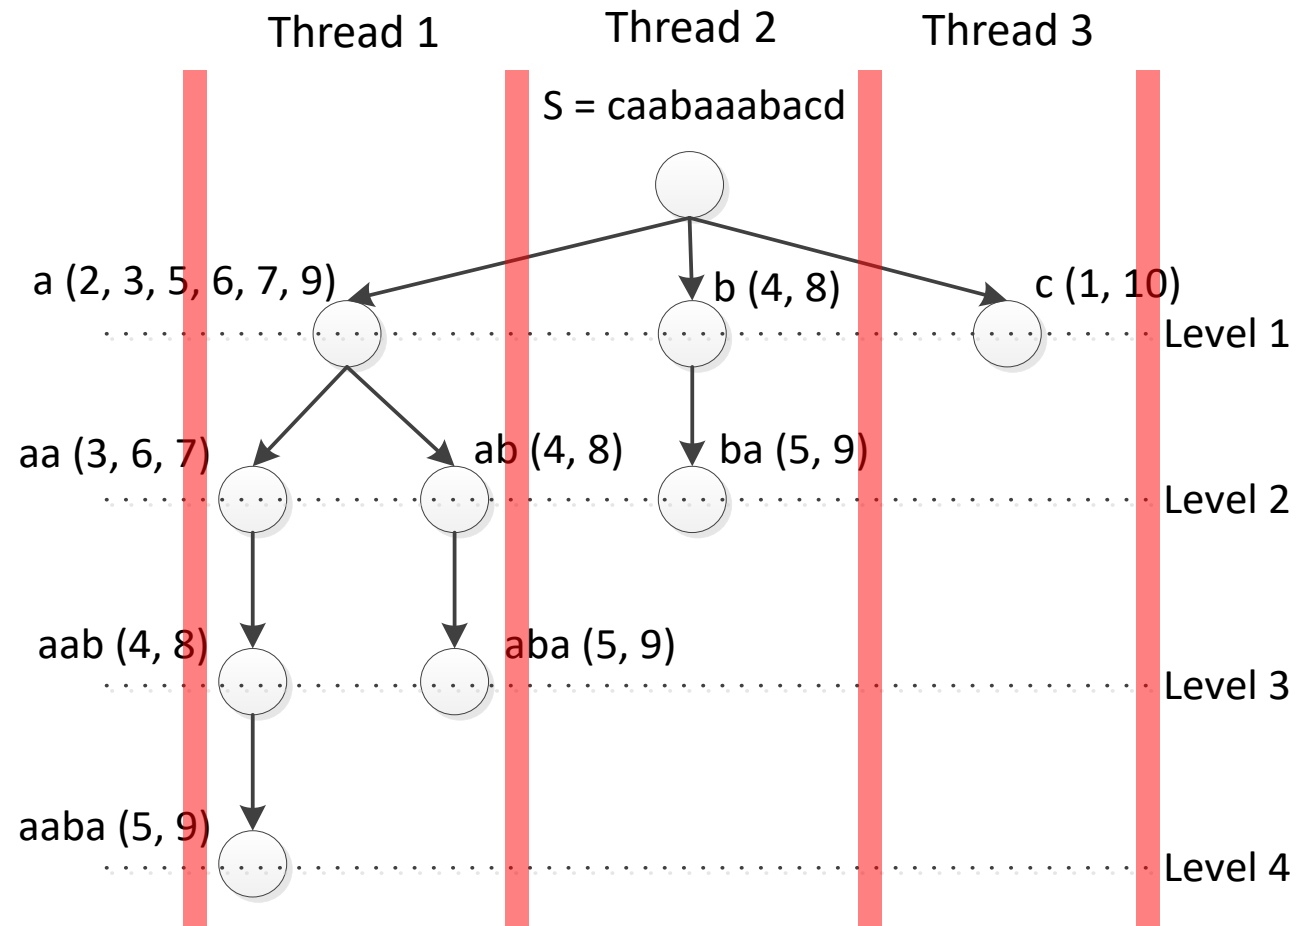
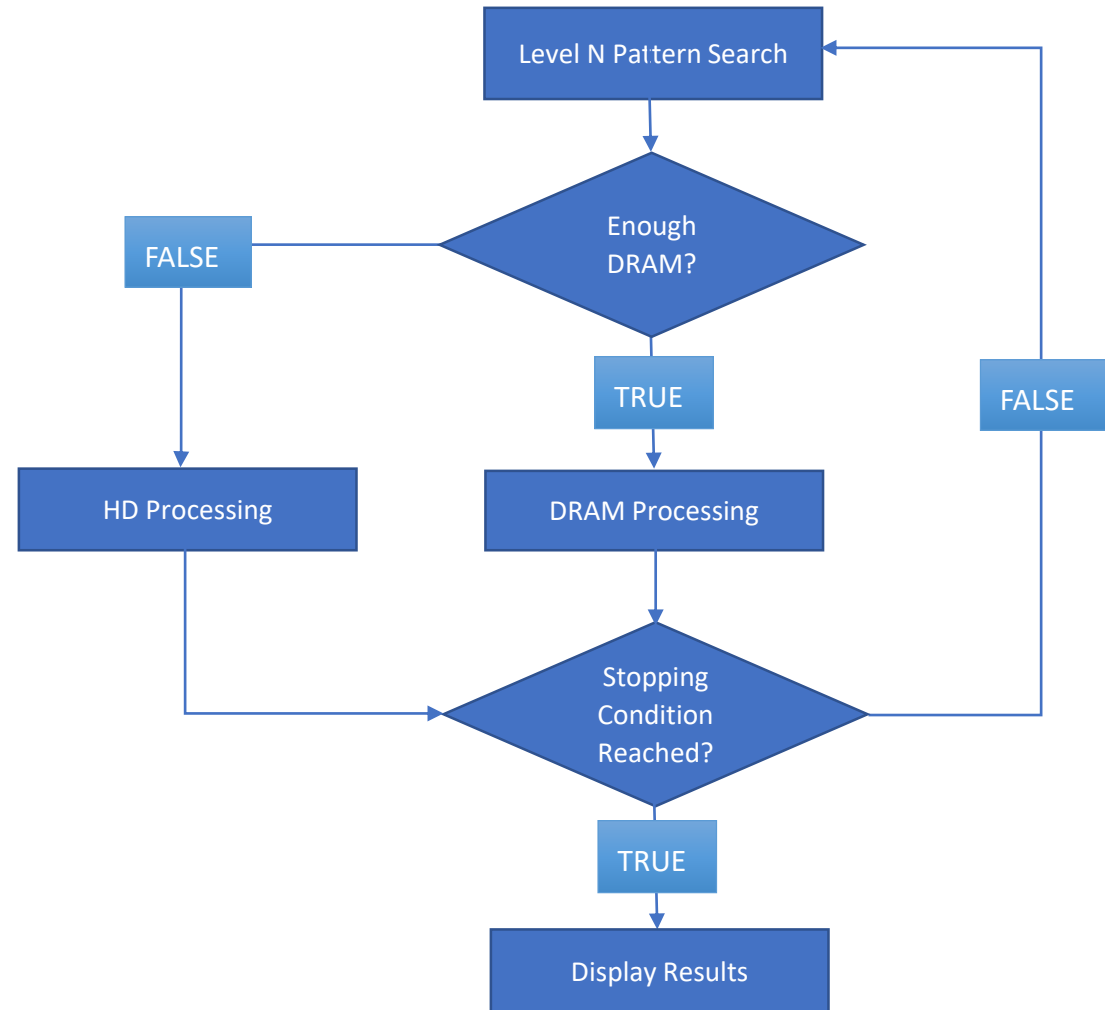


Figure 1: Discovering the subsequences of sequence  $S \rightarrow \text{caabaaabacd}$  having at least 2 occurrences in  $S$ .

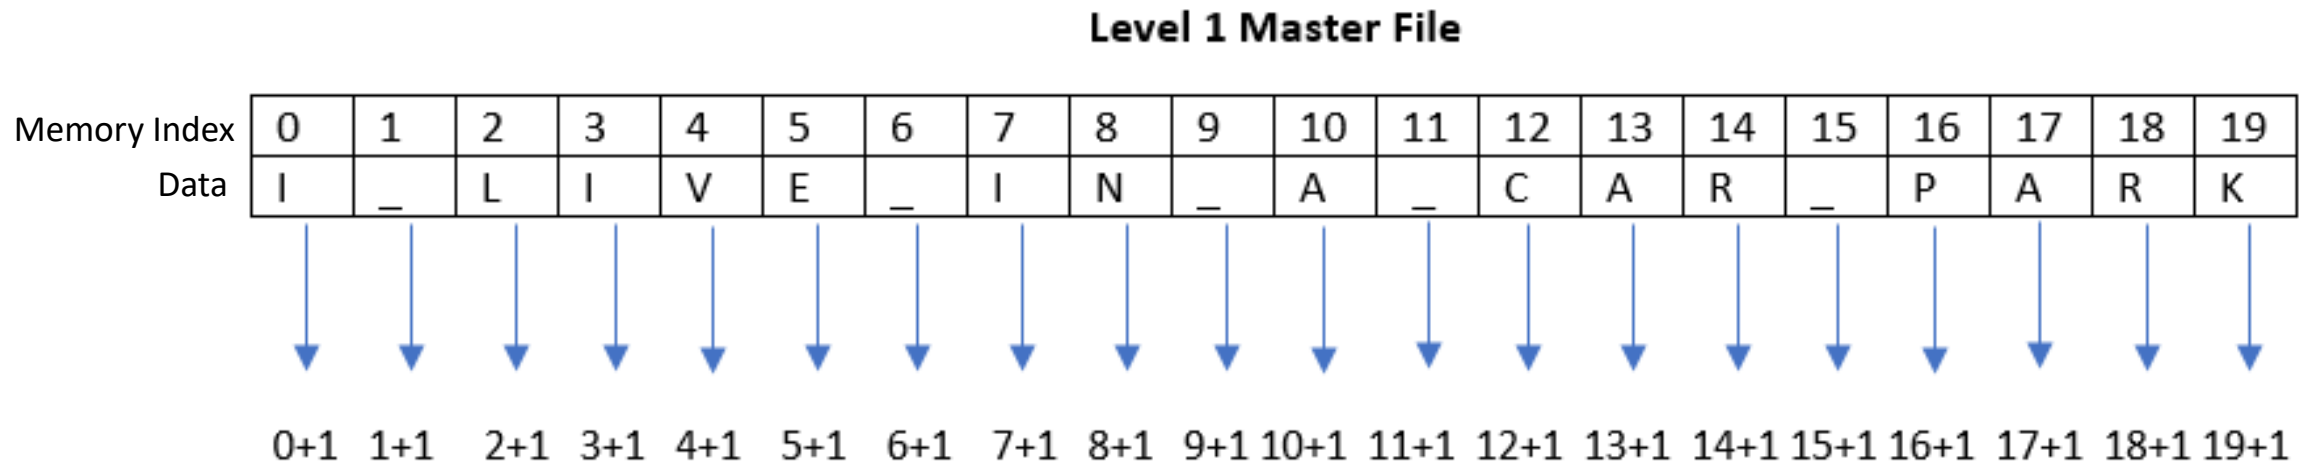
# Pattern Finder Design and Methods

- Core Components of Pattern Finder
  - HD (Hard Disk) Processing
  - DRAM (Dynamic Random Access Memory) Processing
- Pattern Finding using Distributed Computing
- Performance Analysis of Pattern Finder
  - Overlapping vs Non-overlapping Search
  - Multiprocessor scalability
  - HD vs DRAM processing speed comparison

# High Level Processing Flow




# Level Node Processing



- A potential pattern candidate only requires the endpoint position of the pattern sequence
- A pattern can be recovered knowing the level (length) and the last index in the master file


# Level Node Ordering



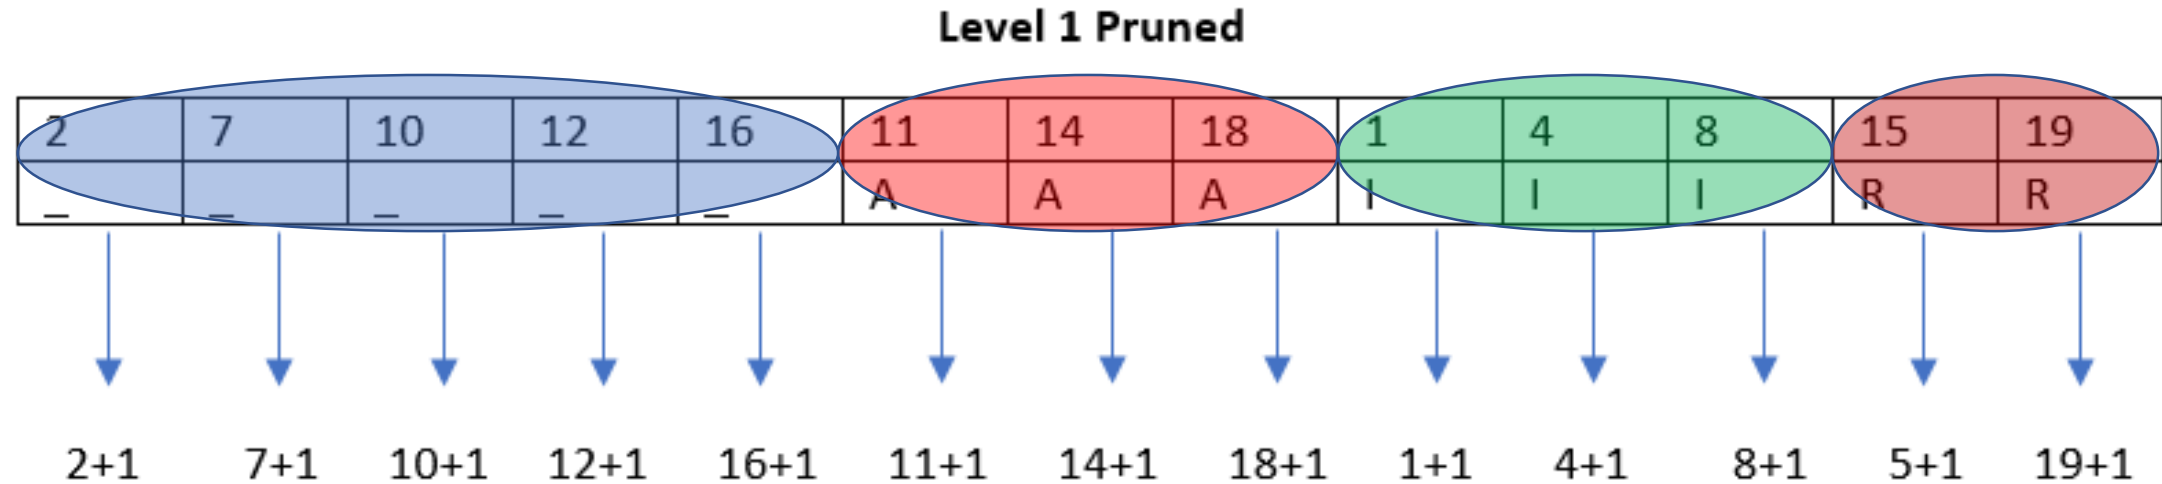
- Each colored vector represents a node in the tree
- Pattern data is ordered based upon ascii mapping 0-255
  - For example A and C in ascii respectively have the values of 41 and 43 and are ordered as such in the array
- Colored circles represent each level node



DISCARDED NODES




# Level Node Pruning




- If a pattern occurs only once, it is pruned
- The next level of pattern searching can be started
- The previous level pattern data can be discarded after the new level data is generated, but for the first level there is no previous data

# Level Searching Resolution




**Level 2 Pruned**



15+1    19+1

**Level 3 Ordered**



**Level 3 Pruned**

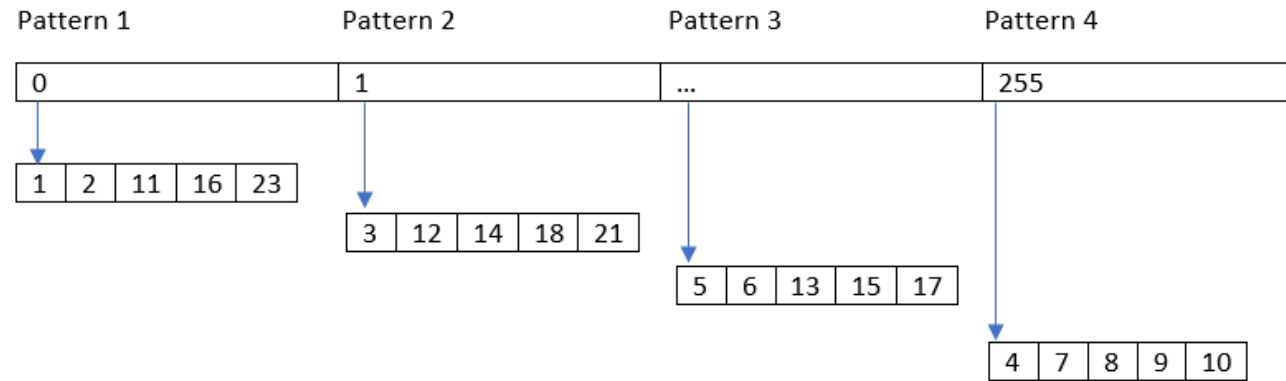
NO MORE PATTERNS!

# Implementation Improvement

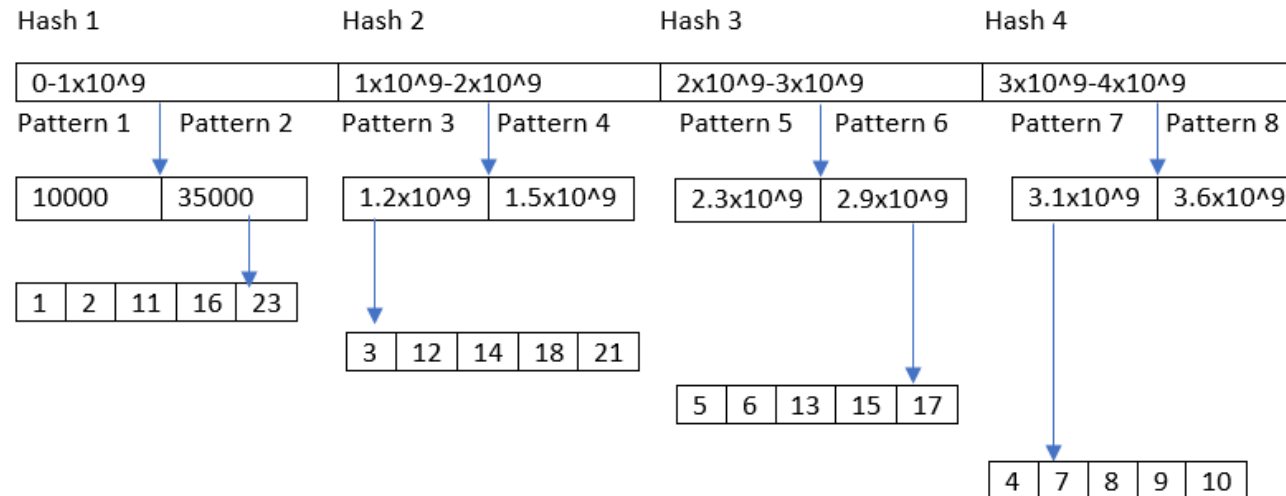
- Previous Pattern Finder algorithm organized next level patterns using a hash map because the integer patterns generated could range between values of 0 - 4294967296.
- New Pattern Finder searches at the byte granularity thus new patterns could be directly mapped to a vector of size 256 that would not require any hash map searches for pattern categorization.
- Searching through large hash maps is very time consuming and by eliminating the need to do map searches, the new implementation runs much faster.

# Direct Indexing Versus Mapping Patterns

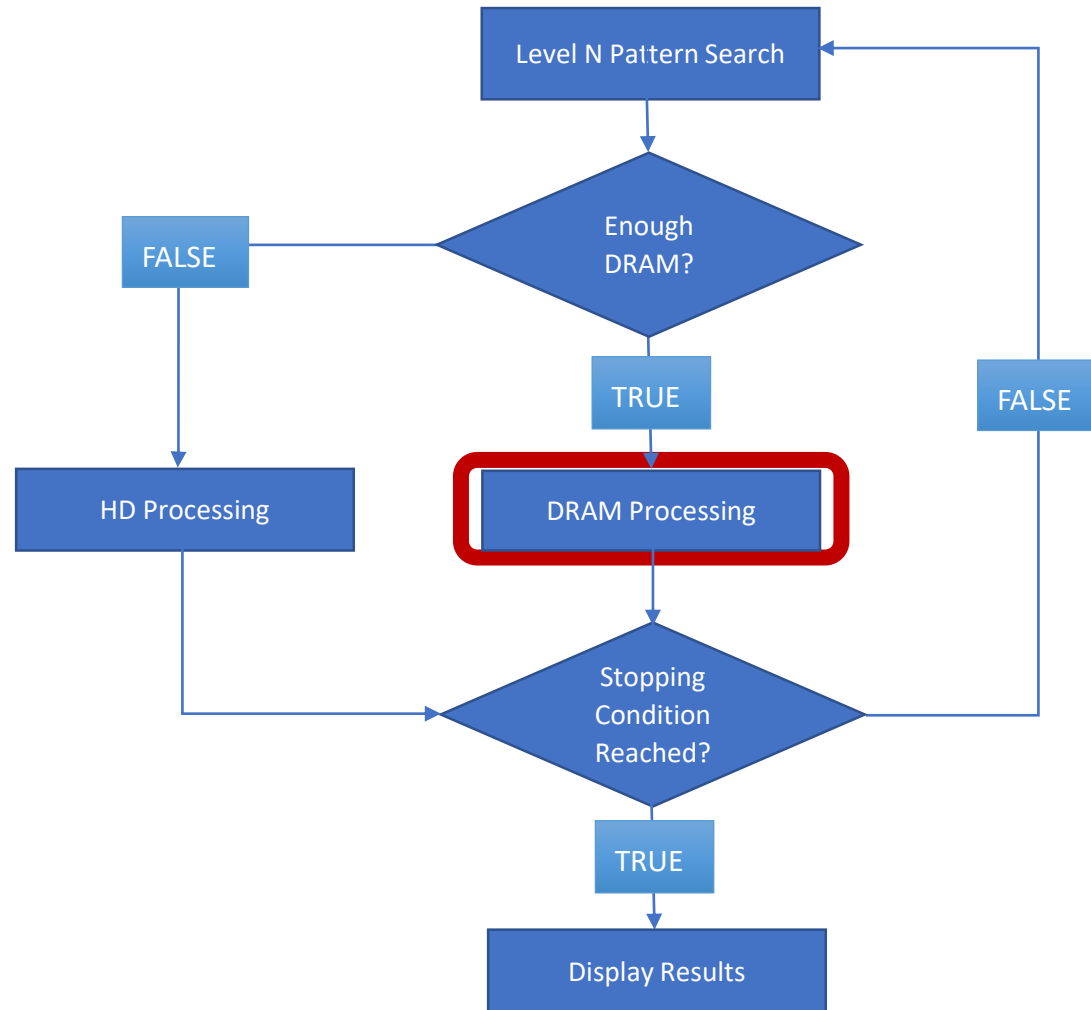
Direct Index Pattern Tracking



Hash Map Pattern Tracking

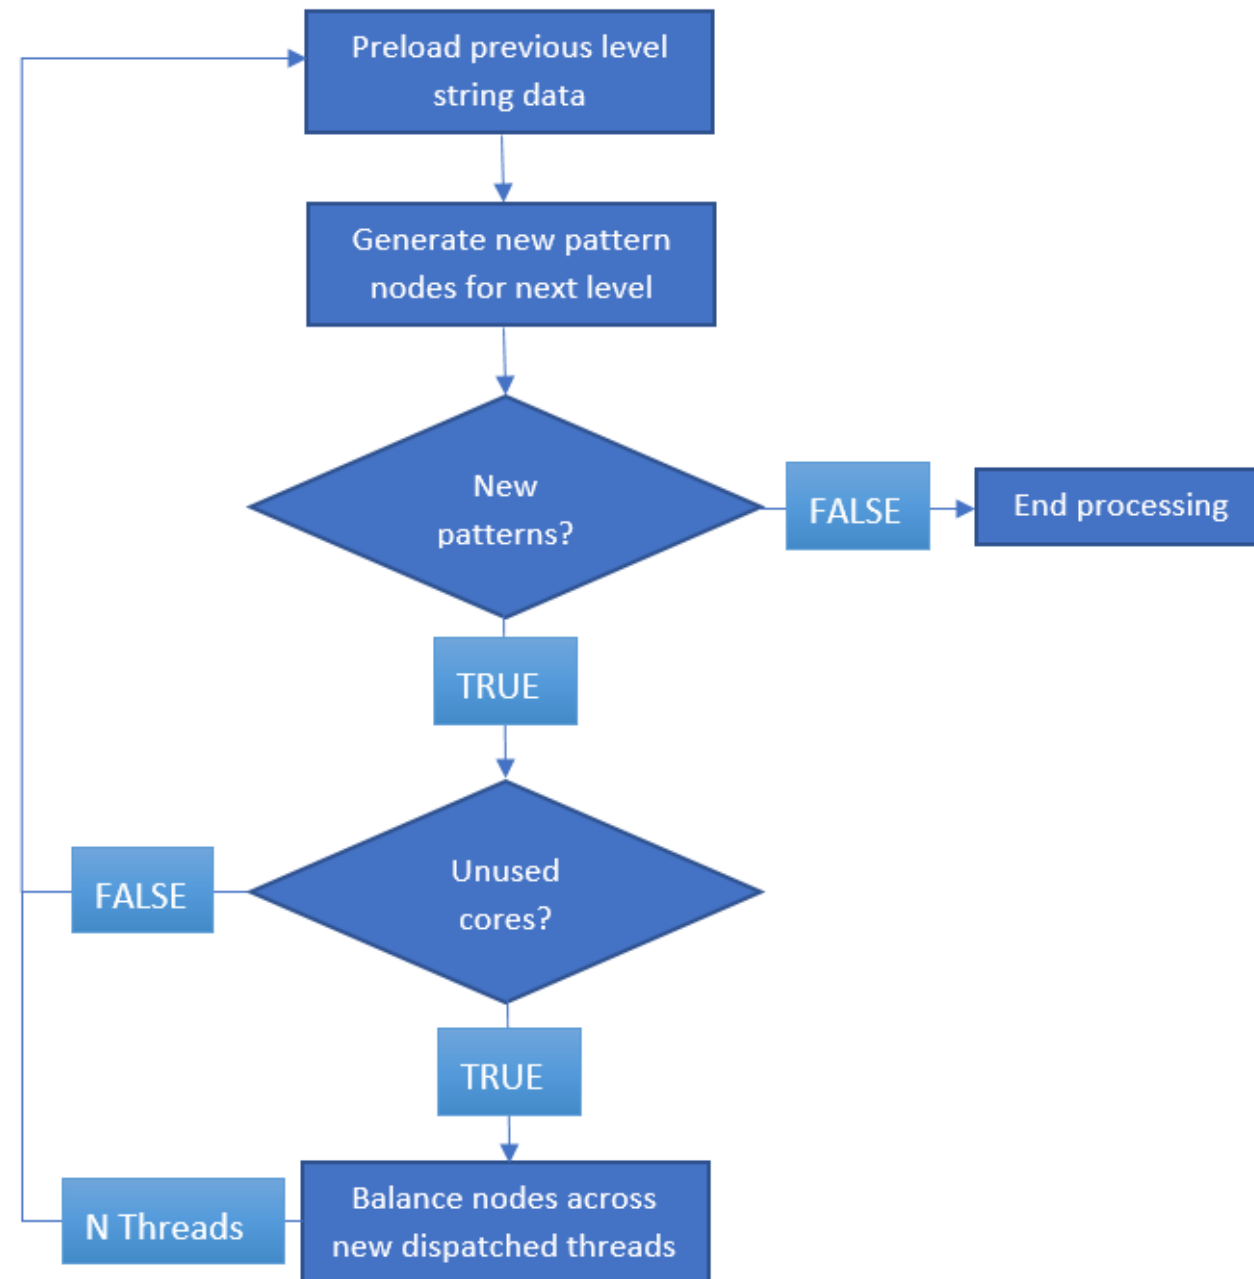


# High Level Processing Flow

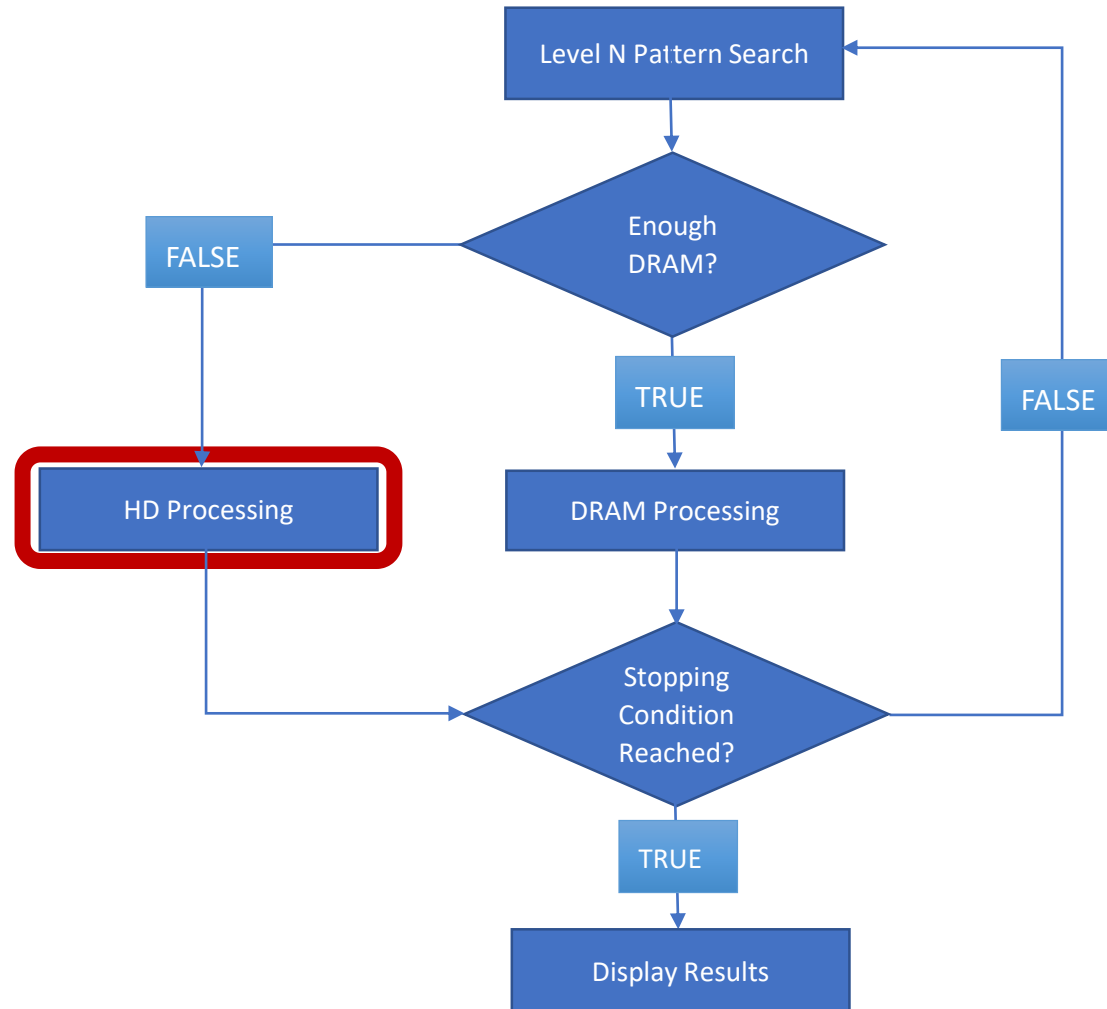


# DRAM Processing

- Stream of reads improves cache locality
- Process preloaded string data and generate new pattern nodes
- If new patterns are generated continue processing
- If unused cores are available then dispatch processing threads recursively which implements thread pooling
- Nodes are distributed evenly across newly spawned worker threads

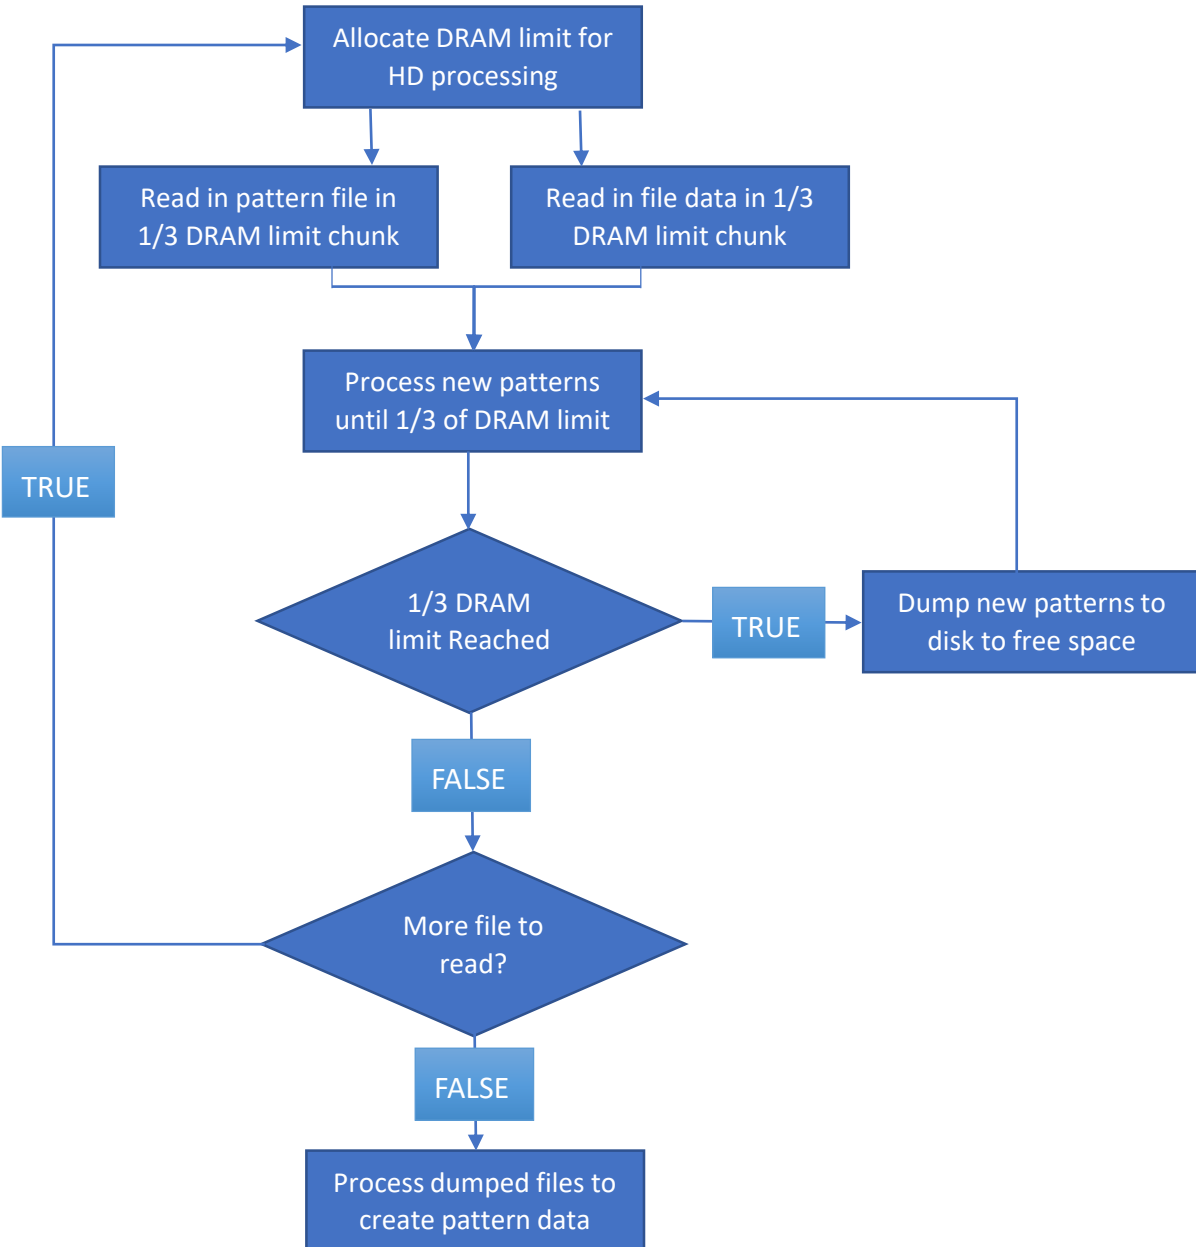


# High Level Processing Flow



# HD Processing


- Allocates three partitions of DRAM space
- The first partition is used to load in a slice of the previous pattern data
- The second partition is used to load in a chunk of the file being processed
- The third partition acts as available memory for new patterns to be created
- When the available memory is completely used it must offload the data to the HD






# File Block Handling

- The red block of memory indicates the loading of file data into DRAM
- Once that allocation size has been reached, the program will not continue to pull the entire file up to DRAM




# Previous Level Pattern Block Handling

- The green block indicates memory allocation of the previous node level data
- If a file containing previous node data is larger than the DRAM allotment then the program will discontinue loading data
- The last loaded node must have the entire pattern index sequence so there typically occurs a negligible overflow
- Not possible to load in segments of a pattern node's index sequence




# Next Level Pattern Block Handling

- When the next level pattern allocations have exceeded the memory allowance the data must be pushed to the hard drive
- Writing the data to the hard drive while deallocating DRAM alleviates space for continued processing

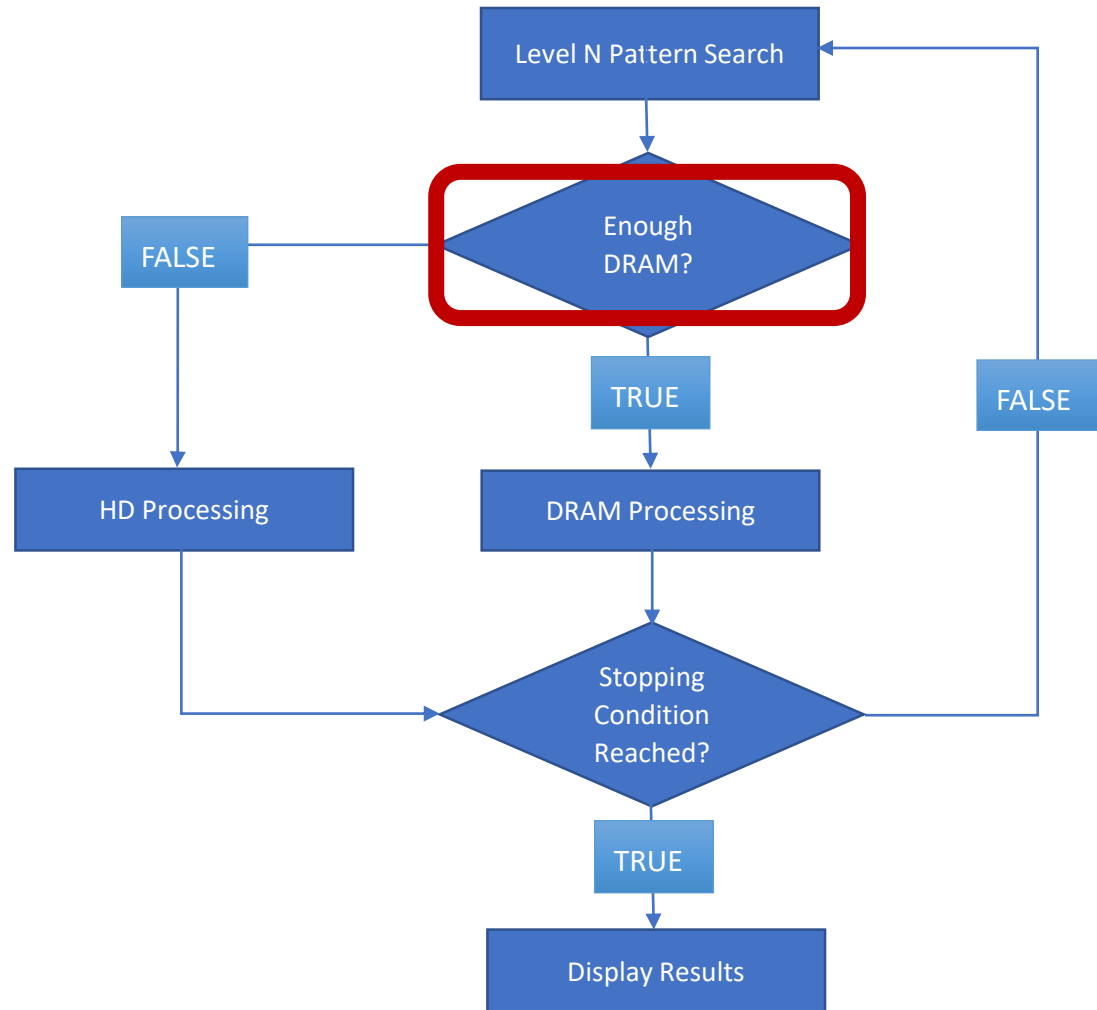


# Offloading Node Data to Hard Disk

- Pattern data is offloaded to the hard disk for later retrieval
- DRAM space is again open for creating new pattern nodes

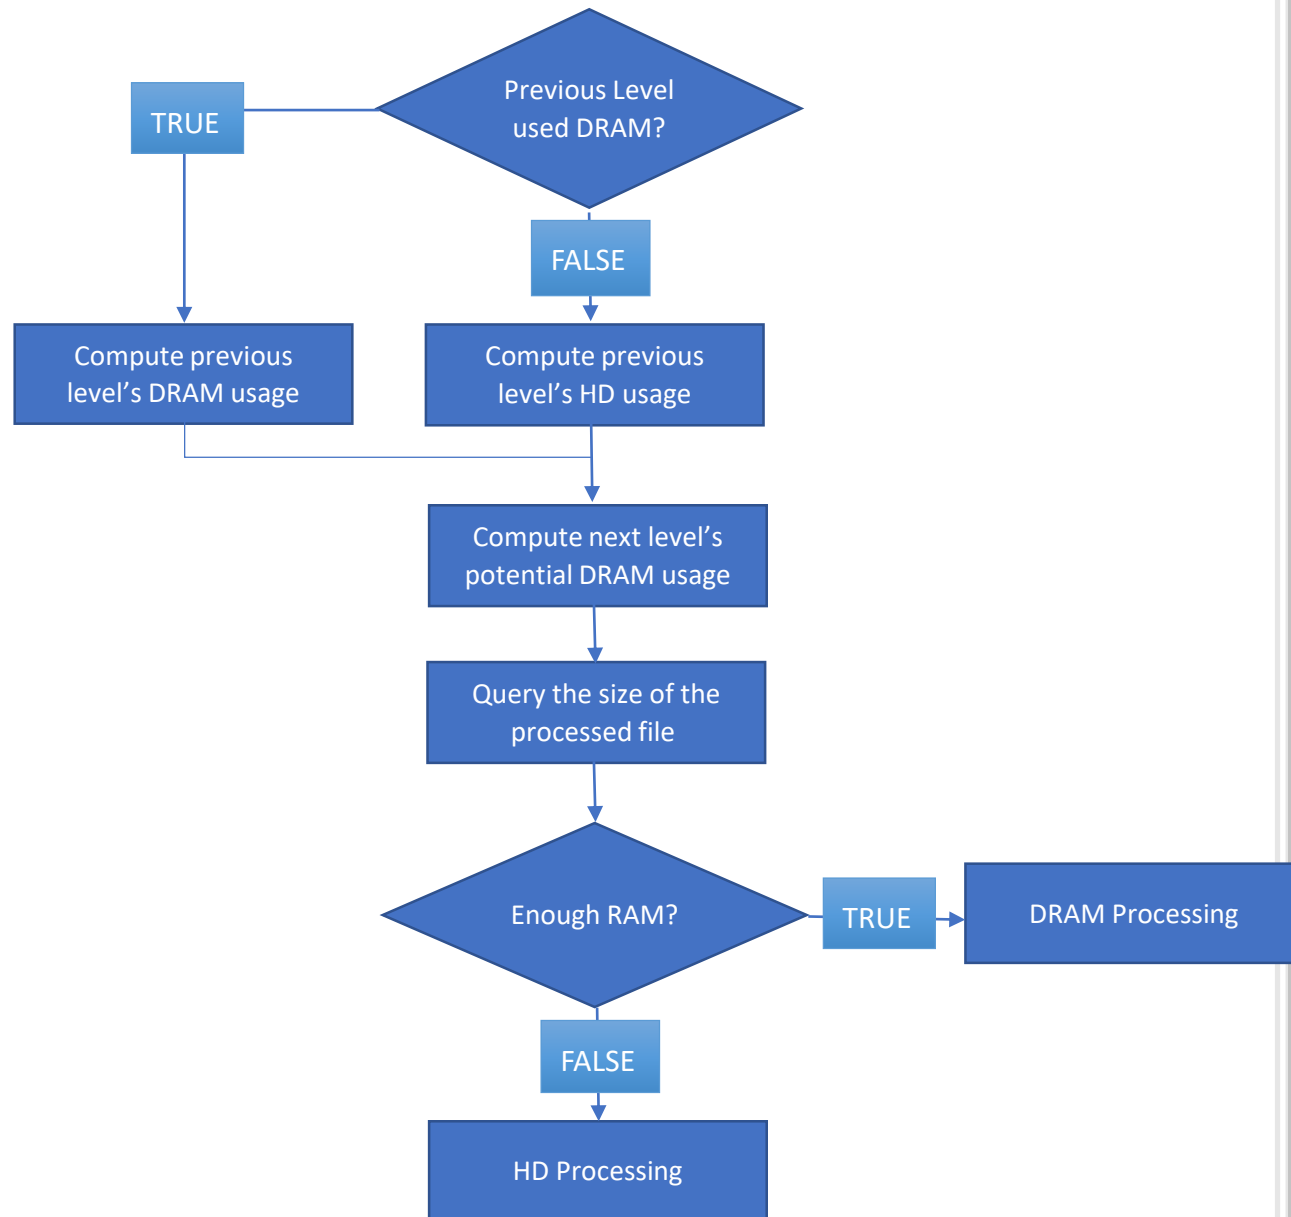


# High Level Processing Flow



# DRAM or HD Processing Decision

- Has the previous level used DRAM or HD? Transitioning between processing in DRAM to HD or vice versa is very time consuming.
- Compute previous level's memory consumption
- Figure out potential patterns generated based on previous pattern data
- Query the size of the file being processed
- After all memory allocations have been added then the decision is made to process with DRAM or DRAM/HD.



# Distributed Computing using Python

- Segment Processing

- 100% accurate, slower but takes N times more memory
- Dispatches N Pattern Finder programs equipped with T threads each to detect patterns that start with a specific byte value.  $N \leq 256$
- Ex. `python segmentedRootProcessing.py ../Database/Data/Boosh.avi N T`
  - Dispatches N processes equipped with T threads each

- Split Processing

- Less than 100% accurate but much faster
- Splits a file into portions and dispatches N single threaded Pattern Finder programs to process the portions individually
- Ex. `python splitFileForProcessing.py ../Database/Data/Boosh.avi N`

# Segment Processing Job Distribution

- Job distribution is key to getting better throughput
- For example a segment job dispatched 8 Pattern Finders each equipped with 4 processing threads which equates to 32 threads of total processing power
- Results
  - 1.5x speedup compared to single Pattern Finder process with 32 threads
  - 4 threads per job is the sweet spot for processor distribution





# Pattern Finder Analysis


- Measure multithreading scalability
- Large data set characterizations
- Memory limited data processing speed comparisons
- Pattern coverage of overlapping and non-overlapping searches
  - Accuracy vs speed tradeoff
- Memory Access Latency statistics
- LLC (Last Level Cache) miss counts

# Multiprocessor Scalability

- The previous scalability results match the 98% parallel program trend line.
- Pattern Finder using 32 threads in parallel results in a 16x speed improvement



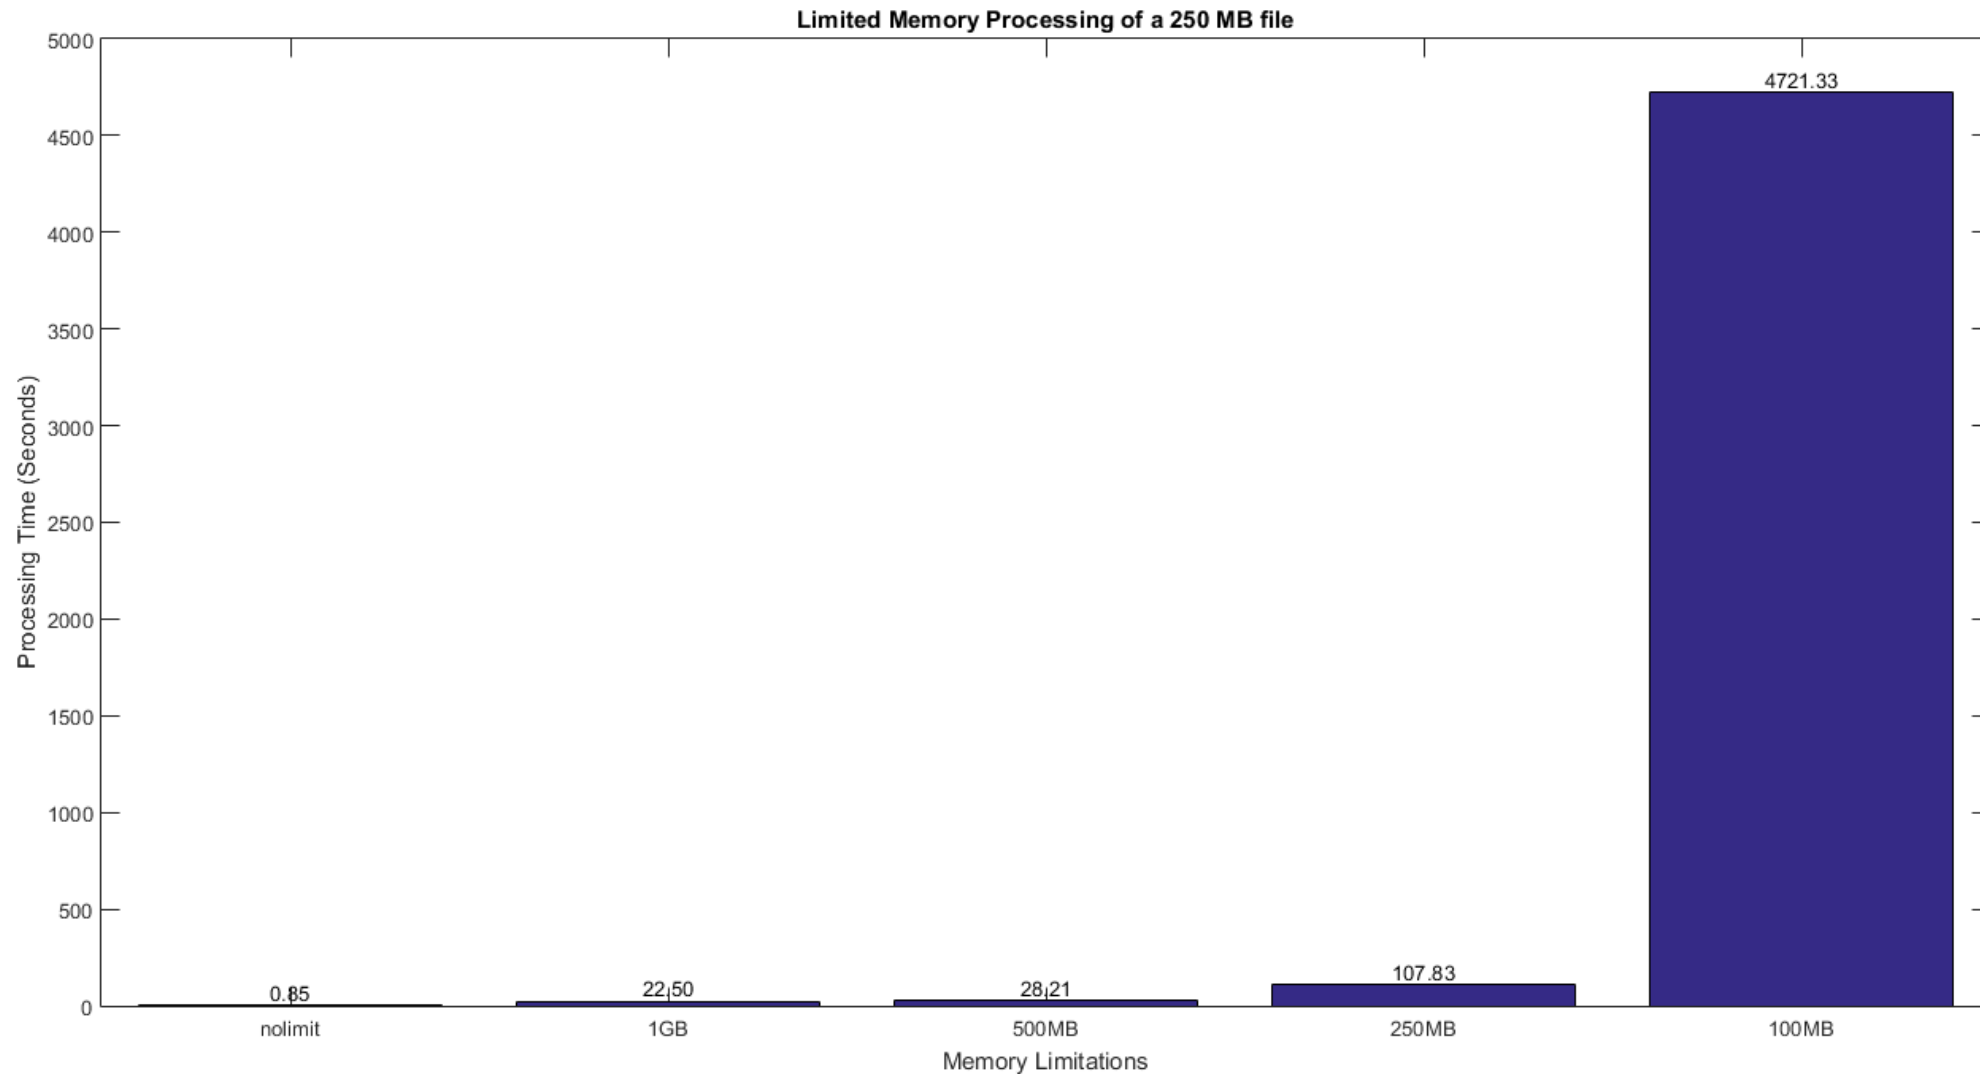
# Dataset Characterizations to Isolate Large Pattern Files



- Large dataset processing identifies pattern characterizations for certain file types
- Deviations from this characterization are isolates files that contain large patterns
- The circled scatter points below the trend line indicate large pattern conditions within in a file
- Files with long patterns take longer to process


# Limited Memory Processing

Memory Limit	Processing Slowdown
100 MB	5554
500 MB	33



# HD vs DRAM Processing

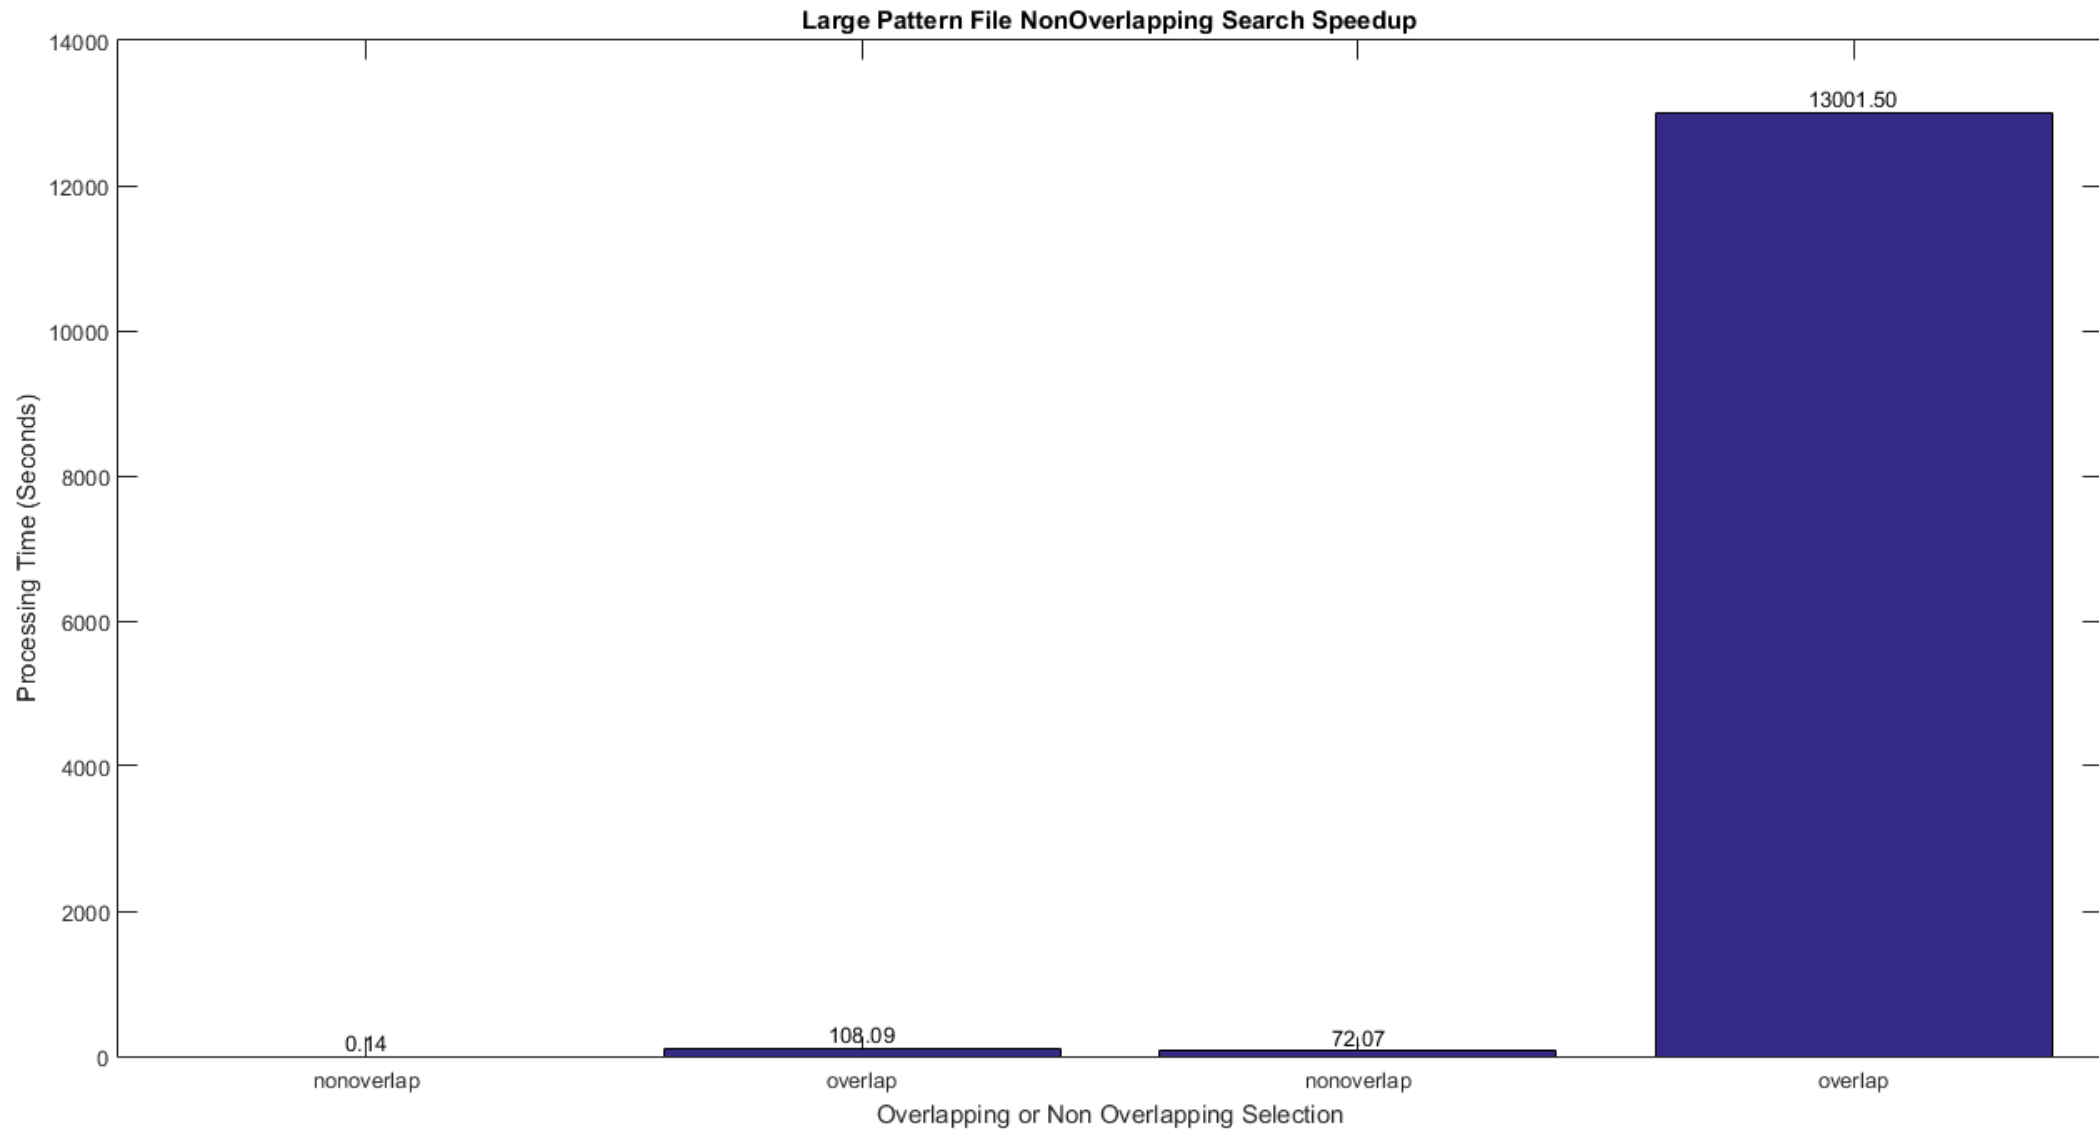
- The HD processing does not scale as well as the DRAM processing as the number of threads increase
- HD processing is about 63 times slower than serial DRAM
- As the threads increase, HD processing starts to linearly slow down at a 2.5x rate per thread compared to DRAM processing.
- HD processing does not scale well due to bad parallelization from reading and writing to the Hard Disk.



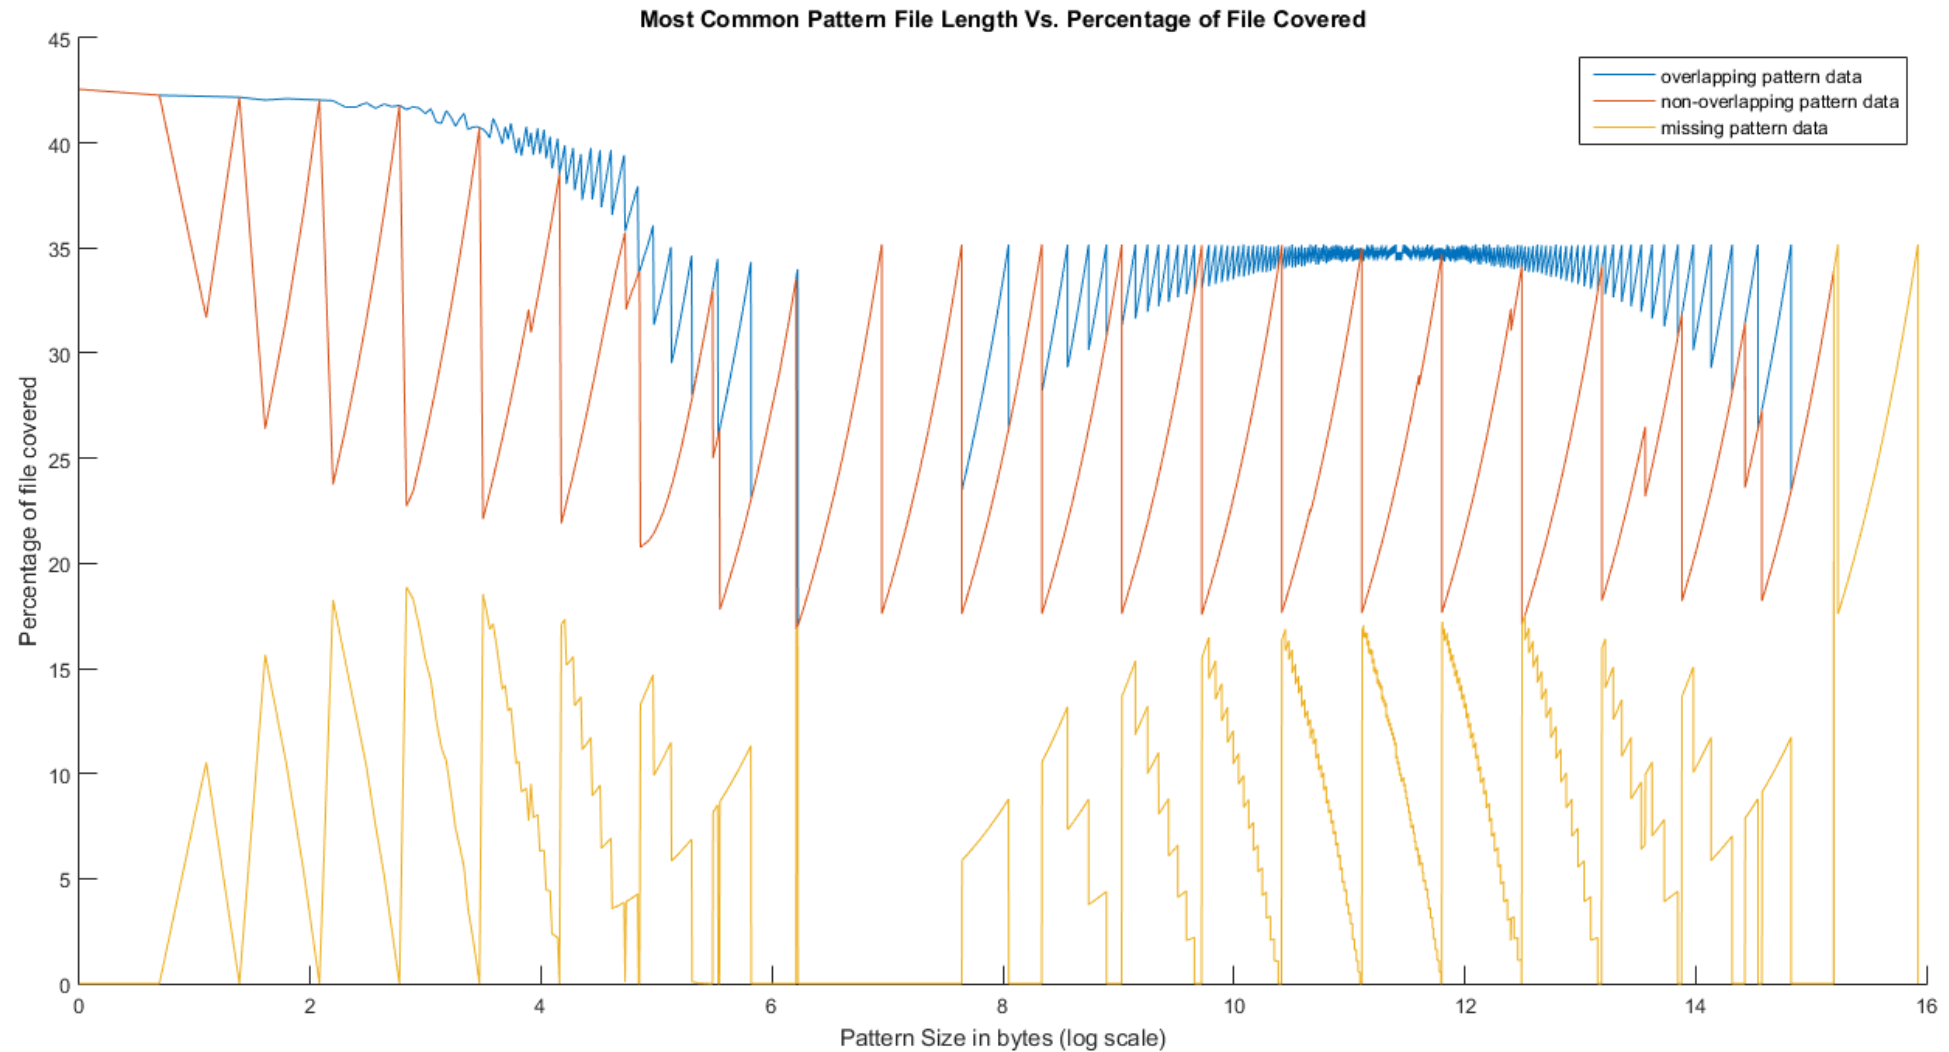
# Overlapping vs Non-overlapping Search

- Files containing large patterns benefit greatly in speed when utilizing the non-overlapping search.
- The processing time is much faster while maintaining nearly 100% pattern accuracy
- Instances within a pattern set can be removed for overlapping while pattern instances that overlap other pattern instances do not get pruned.

# Files Containing Large Patterns Speedups




# Coverage of Files Containing Large Patterns







# Non-overlapping Pattern Resolution



Patterns of length 1 with 100% coverage




Patterns of length 2



Resolution of Overlap Patterns of length 2 with 100% coverage



Patterns of length 3



Resolution of Patterns of length 3 with 60% coverage



# Profiling and Optimizing Pattern Finder

- Linux tools
  - Perf used for finding program bottlenecks
    - Performance test commands
      - Perf record -g ./PatternFinder -f Boosh.avi -v 1 -threads 64 -ram
      - Perf report -call-graph
    - Valgrind used for memory leak detection
  - Windows tools
    - Intel Profiling Suite is a powerful memory leak and cache metric analyzer
    - Visual Studio profiler used for bottleneck detection
  - Hardware Platform
    - 4 Xeon processors totaling 96 cores of processing power with 1 TB of RAM


# Memory Bottleneck Results

<b>Elapsed Time</b> <sup>?</sup> :	<b>85.792s</b>
<u>CPU Time</u> <sup>?</sup> :	615.561s
✓ <u>Memory Bound</u> <sup>?</sup> :	<b>50.0%</b>
<u>L1 Bound</u> <sup>?</sup> :	0.060
<u>L2 Bound</u> <sup>?</sup> :	0.013
<u>L3 Bound</u> <sup>?</sup> :	0.052
<u>DRAM Bound</u> <sup>?</sup> :	<b>0.304</b>
<u>Loads</u> :	733,572,200,710
<u>Stores</u> :	299,352,490,220
<u>LLC Miss Count</u> <sup>?</sup> :	5,718,343,080
<u>Average Latency (cycles)</u> <sup>?</sup> :	10
<u>Total Thread Count</u> :	21
<u>Paused Time</u> <sup>?</sup> :	0s

- 250 MB video file
- Cache management is essential to improving speed when scaling up with threads
- The average latency for data access is 10 cycles for this file which corresponds to an L2 cache hit performance metric
- Latency is dependent upon file composition so results may differ

# Cache Spatial Locality in Large Files

- Most files processed will be very large and therefore create memory jumps in a node's address list when a node transition occurs
- Within the T node there are memory transitions from 4 to 70 which will not be covered in one 64 byte aligned cache block
- The memory transition from T to U is 7050 to 7 which is another cache miss
- All these memory hops most likely give the L2 cache latency metric



# Pattern Finder as a Tool

- Generates most common pattern data per level
- Generates pattern data for overlapping and non-overlapping searches for post processing coverage comparisons
- Compounds pattern data for large file databases
- Multiple instances of Pattern Finder can be dispatched for distributed computing solutions.

# Pattern Finder Options

- -f [filename | filedirectory] specific file or directory to pattern search
- -n indicates a non overlapping pattern search
- -i [patterncount] number of instances of a sequence that constitutes a pattern, default is 2
- -min [patternsize] indicates first pattern size to begin searching for
- -max [patternsize] indicates the largest pattern size that can be searched
- -threads [threadcount] indicates threads to use for processing
- -mem [MegaBytes] indicates memory limit for processing

# Pattern Finder Examples

- `./PatternFinder -f Database -threads 64 -ram`
  - Pattern searches all files recursively in directory using DRAM with 64 threads
- `./PatternFinder -f TaleOfTwoCities.txt -min 5 -max 100`
  - Finds patterns of length 5 to 100 and then terminates processing
- `./PatternFinder -f Boosh.avi -n`
  - Processes file using non overlapping processing
- `./PatternFinder -f StairwayToHeaven.mp3 -mem 1000`
  - Processes file using memory prediction per level for HD or DRAM processing with a memory constraint of 1 GB

# Summary and Future Work

- Thesis Accomplishments

- A multi threaded Pattern Finder resulted in up to 23 times faster processing.
- The best example of improvement stems from a trace file that took 12.5 days to process single threaded, 36 hours to process multithreaded and 3 hours to process using a non overlapping multithreaded search.
- Processing large files using limited memory is possible but the processing times are still very slow if the memory limit is too small.

- Future Work

- Find a better way to limit a program's DRAM usage
- Implement a job distributed pattern finder script to process files on a super computer grid.



# Acknowledgements

- Resit Sendag for mentoring me and guiding me through the implementation of a multi threaded Pattern Finder
- My parents for keeping me sane